

Agent-Population Scaling and Decentralized Team Recruitment in Alem

Combined E0–E3 Methods, Results, and Preregistered Evaluation Report

AlemDICE Research Engineering Report

23 July 2026

Abstract

This report consolidates a staged investigation of physical-agent scaling, decentralized team recruitment, and team-scoped communication in Alem. The completed evidence is limited to two provider-free engineering studies. E0 exercised the canonical evaluator at $N \in \{1, 2, 3, 4, 6\}$ for two seeds and ten ticks and passed every structural, accounting, replay, and renderer gate. E0.1 requested 200 ticks with the same scripted Noop policy. All structural gates again passed, but only seven of ten episodes reached the requested horizon; three ended normally after all agents died. E0.1 is therefore a structural **PASS** and a literal full-horizon **FAIL**. Neither study measures useful behavior.

E1, the E2b v3 language-model screen, and E3 remain prospective. E2b v2 made a bounded hosted attempt, but its one-method canary failed to expose a Mutual-Nomination output-limit failure; the promoted stage is preserved as a poisoned, non-promotable diagnostic rather than an efficacy result. A provider-free E2d2 replication promoted a deterministic joint exact allocator under truthful public claims; this is allocation evidence rather than an LLM or Alem efficacy result. E1 will measure the unchanged GPT-5.4 nano Source baseline as physical population changes. E2 isolates decentralized recruitment in a deterministic six-agent microbenchmark with bounded public control messages and team-private conversation. E3 will transfer a promoted mechanism to Alem behind an opt-in topology while preserving Source behavior. The report fixes hypotheses, paired seeds, outcomes, promotion rules, audit requirements, and implementation risks before those results are observed. In particular, population scaling is treated as a fixed-world descriptive curve rather than a pure causal agent-count effect, and no claim about decentralized recruitment will be made unless routing, agreement, validity, and replay gates pass.

1 Program status and claim boundary

The program asks three ordered questions:

1. **Population scaling:** How does the unchanged Source language agent behave as the number of physical agents changes?
2. **Formation mechanism:** Can six embodied agents form useful, task-bound teams using bounded decentralized recruitment records?
3. **Environment transfer:** Does a promoted recruitment mechanism retain useful performance in Alem when ordinary conversation is restricted to current teammates?

Each stage removes a distinct ambiguity. E0 separates infrastructure from model behavior. E1 measures population behavior before routing is altered. E2 separates recruitment from the full game. E3 tests whether the mechanism survives the navigation, survival, and action constraints of Alem. A failure at an earlier gate blocks downstream efficacy claims.

The completed sequence may support claims about population scaling, decentralized team formation, and team-scoped communication. It does not by itself establish strict information isolation, robustness to adversarial or crashed agents, learned communication, long-horizon role alignment, or populations above six.

Stage	Status	Current evidence or blocker
E0: 10-tick compatibility	PASS	Ten provider-free episodes completed; all structural gates and five renderer smokes passed.
E0.1: 200-tick extension	PASS structural / FAIL horizon	Ten complete artifact sets; 7/10 episodes reached tick 200; three terminated naturally after all-agent death.
E1: Source population curve	PENDING	Profile, logging, launcher, and analysis are implemented; the provider preflight and behavioral evaluation have not yet run.
E2: RecruitmentArena-6	PASS provider-free / FAIL v2 / PENDING v3	TFP1, arena, policies, selectors, oracle, and the staged E2b harness are implemented. E2d2 promoted joint exact allocation. E2b v2 failed after 72 hosted calls; v3 is prospectively amended and has passed fake orchestration only.
E3: Alem recruitment	PENDING	Blocked on E2 promotion and a correct opt-in, receiver-specific communication route.

Table 1: Status at the report date. “Pending” denotes an unobserved study, not a null or negative result.

2 Shared experimental conventions

2.1 Behavioral model and unchanged Source condition

All paid behavioral stages use homogeneous GPT-5.4 nano action agents with reasoning effort **high**. The Source arm retains:

- the `robust_all` agent;
- the `specific_collaborative` prompt;
- normal reasoning and scratchpad memory;
- the existing action parser and legal-action affordances;
- one model call per physical agent per environment tick; and
- ordinary one-tick-delayed peer broadcast.

No bodyless leader, extra planner, or additional planning call is part of this program. Agent calls within a tick run concurrently. Behavioral arms use one episode worker unless a separate provider-load experiment is preregistered.

2.2 Paired designs and terminal episodes

Comparisons use identical environment seeds across arms. A natural all-agent death is a valid terminal outcome, not an API failure. Every report must state both the requested horizon and actual executed ticks. Cost denominators include actual physical agent-ticks so early death cannot appear to improve efficiency.

2.3 Primary outcomes

The three co-primary Alem outcomes are:

1. paper Total achievement percentage;
2. mean per-agent episode return; and
3. raw unique team achievements.

Supporting outcomes include Base and Coordination achievement percentages, deaths, alive-agent steps, action parse rate, task attempts and successes, messages, emitted bytes, delivered bytes, tokens, model latency, and complete tick wall time. Formation studies additionally report normalized task reward, oracle regret, feasible-team rate, formation latency, churn, roster agreement, and truthful-report calibration.

3 E0: short variable-agent compatibility gate

3.1 Method

E0 crossed $N \in \{1, 2, 3, 4, 6\}$ with seeds 13000 and 13001 for ten ticks in `alem/default`, Easy coordination difficulty, and baseline communication. Exactly one indexed client configuration existed per physical agent, but a deterministic test actor replaced inference. Every actor selected `Noop` and emitted

$$\text{E0} \mid \text{AGENT} = i \mid \text{TICK} = t \mid \text{STATE} = \text{ACTIVE}.$$

The normal evaluator still performed reset, stepping, action canonicalization, one-tick routing, trajectory and state serialization, debug output, terminal accounting, and attempt-ledger writing.

The gate checked:

- client and agent cardinality and every trajectory agent axis;
- Warrior, Forager, Miner specialization cycling;
- every ordered `Give to Agent  $j$` mapping;
- action parse rate equal to one;
- zero model calls, provider requests, and transport errors;
- exact baseline accounting, $M_{\text{emit}} = NT$ and $M_{\text{deliver}} = N(N - 1)T$;
- complete canonical artifacts and ten state snapshots per episode; and
- equality of stable reset and saved-state projection hashes.

One seed-13000 state at every population was also encoded as a 768-by-832 H.264 full-world renderer smoke artifact and validated with `ffprobe`.

3.2 Measured result

N	Episodes	Ticks	Observation dim.	Actions	Give maps	Emitted	Delivered	Delivery bytes	Result
1	2	20	9,476	55	0	20	0	0	PASS
2	2	20	9,597	55	2	40	40	1,200	PASS
3	2	20	9,718	56	6	60	120	3,600	PASS
4	2	20	9,839	57	12	80	240	7,200	PASS
6	2	20	10,081	59	30	120	600	18,000	PASS

Table 2: E0 measured results. Give mappings count ordered giver–recipient pairs, hence $N(N - 1)$.

Across ten episodes, E0 executed 100 environment ticks and 320 scripted agent-ticks, emitted 320 messages, delivered 1,000 copies totaling 30,000 bytes, and saved 100 states. All structural and renderer gates passed. Observation length increased by 121 values for each added teammate. Baseline delivery fan-out grew quadratically. These are interface and accounting facts, not behavioral benefits.

3.3 E0 decision

E0 removed the short-run engineering blocker for the supported $N = 1, 2, 3, 4$ screen and for the separately labeled $N = 6$ extension. It did not test task performance, useful coordination, model reasoning, provider concurrency, recruitment, or role coherence.

4 E0.1: 200-requested-tick compatibility extension

4.1 Method and decision rule

E0.1 repeated the same populations, seeds, scripted actor, Easy environment, and baseline topology with 200 requested ticks. It added monotonic timing for the full episode, the complete evaluator tick, the slowest tick, and the concurrent scripted-worker phase.

Two decisions were fixed:

Structural gate.

Natural termination is allowed, but every episode must produce complete artifacts and pass all shape, action, role, communication, replay, and provider-accounting checks.

Strict horizon gate.

Every episode must execute all 200 requested ticks in addition to passing the structural gate.

4.2 Measured gate outcomes

All ten artifact sets were complete and every structural check passed. Seven episodes reached tick 200. Three ended normally when all agents died from mob combat:

- $N = 1$, seed 13000: 114/200 ticks and one death;
- $N = 4$, seed 13001: 182/200 ticks and four deaths; and
- $N = 6$, seed 13001: 192/200 ticks and six deaths.

Consequently, E0.1 is a structural **PASS** and a strict-horizon **FAIL**. The failure is retained rather than being reclassified after observing that it was a clean environment termination.

N	Actual/requested ticks	Mean return	Full horizon	Alive-step %	Deaths/slots	Mean tick (s)	Mean excl. max (s)	Structural	Horizon
1	314/400	-0.4500	1/2	N/A	1/2	0.151301	0.076475	PASS	FAIL
2	400/400	-0.5250	2/2	95.88	1/4	0.175483	0.124436	PASS	PASS
3	400/400	-0.0333	2/2	100.00	0/6	0.235863	0.179381	PASS	PASS
4	382/400	-0.8625	1/2	78.73	7/8	0.296600	0.236695	PASS	FAIL
6	392/400	-0.6833	1/2	86.48	8/12	0.437702	0.379037	PASS	FAIL

Table 3: E0.1 structural, terminal, descriptive return, and timing results. Return is mean cumulative reward per physical agent. “Mean excl. max” removes exactly one slowest tick from each episode and is descriptive rather than a separately timed warm benchmark.

The matrix executed 1,888/2,000 requested environment ticks and 6,194/6,400 requested scripted agent-ticks. It recorded 6,194 emitted messages, 19,544 delivered copies, 614,408 delivered bytes, 1,888 saved states, and zero model calls, provider requests, or transport errors. The summed episode wall time was 498.207821 seconds; end-to-end elapsed time was 659.272638 seconds. The step-weighted complete-tick mean was 0.263203 seconds, and the scripted-worker phase totaled only 0.326201 seconds.

Every achievement, normalized task-reward, player-progression, and monster-kill field was zero. Multi-agent coordination and item-give attempt counts were also zero. The nonpositive return therefore reflects net health change under **Noop**, not task accomplishment. The two-seed returns and mortality patterns cannot rank populations.

4.3 E0.1 decision

E0.1 strengthens the engineering evidence that the evaluator and artifact path support variable populations over substantially longer runs. It does not show that an inactive policy survives to a fixed horizon. Behavioral experiments must treat all-agent death as an outcome and normalize exposure by actual agent-ticks.

5 E1 preregistration: Source performance versus population

5.1 Questions and hypotheses

E1 asks how the unchanged Source condition behaves as physical population changes. It is intentionally descriptive because changing N also changes spawn geometry, role balance, mob slots, prompt size, teammate observations, communication fan-out, and some coordination requirements.

H1a: behavioral scaling.

Useful task performance will vary detectably with population, but no monotonic direction is assumed in advance.

H1b: communication scaling.

Under full broadcast, delivered message copies and bytes will increase approximately with $N(N - 1)$ conditional on emission rate.

H1c: efficiency tradeoff.

Any performance gain at larger N may be offset by greater token, communication, and complete-tick cost.

5.2 Fixed configuration and funnel

The behavioral model is GPT-5.4 nano at high reasoning with `robust_all`, `specific_collaborative`, `scratchpad`, `action parser`, and `free peer broadcast`. `Debriefs` and `bulk image generation` are disabled. Agents within a tick are called concurrently; episodes use one worker.

The executable profile is `source_scaling_200`. It declares six client slots but the launcher overrides only physical population size; surplus slots are never instantiated as agents. Provider reasoning is controlled by `reasoning_effort=high`; the local-model `agent.reasoning` switch remains unset. Population-and-agent-specific prompt-cache keys use one stable traffic shard, so an interrupted resume returns to the same route rather than reassigning a remaining seed to another seed's shard.

Before any episode call, a dedicated cache route makes one bounded GPT-5.4-nano-high compatibility request and requires an exact parse of `<action>Noop</action>`. Its source commit, dependency lock, resolved configuration hash, response ID, provider status, usage, latency, transport attempts, and effective cache key are written durably. The episode campaign then has a nominal ceiling of 6,000 logical calls for E1a and 3,600 for E1b. Hard campaign ceilings are 12,001 successful logical responses including the preflight and 15,000 observed provider attempts. An exclusive run lock, atomic manifests, exact cell-identity checks, append-only attempt ledgers, and child-process cleanup protect paid resume.

Each terminal artifact includes the `alem-dice-performance-v1` block. Base, Coordination, and Total are reward-weighted paper scores on a 0–100 scale. Unique team achievement counts are binary first-unlocks by achievement type; summed per-agent counts may count the same type once for each attaining agent. They are not repeated world-event counts. Separate cumulative counters retain coordination attempts, resolved attempts, successes, transfers, requests, and revives without summing heterogeneous event units.

E1a must pass at a supported population before that population enters E1c. E1b is separately labeled as an extension beyond the documented one-to-four range. $N = 1$ is plotted as a solo reference because it uses the single-agent wrapper.

Stage	Populations	Seeds	Horizon	Purpose
E1a	1, 2, 3, 4	13100–13102	200	Supported-range screen
E1b	6	13100–13102	200	Gated extension
E1c	Every passing N	13200–13209	1,000	Descriptive curve
E1d	3 and 6	13300–13319	10,000	Paper-scale confirmation

Table 4: E1 staged design. E1d reports Easy, Medium, and Hard separately and begins only after prior gates pass and an updated cost estimate is archived. The campaign was authorized to proceed without another approval pause on 23 July 2026.

5.3 Analysis and planned outputs

The headline figure contains three panels:

1. paper Total achievement percentage;
2. mean per-agent episode return; and
3. raw unique team achievements.

Every panel shows raw seed points and 95% bootstrap intervals. Supporting panels show Base and Coordination achievement percentages, deaths, alive-agent-step exposure, action parsing, total and per-agent-tick tokens, delivered bytes, and complete-tick latency.

Planned figure: Performance versus physical-agent count, with Total percentage, mean per-agent return, and raw unique team achievements as the three primary panels.

Pending result. Populate this subsection only from canonical E1 artifacts. Report early terminations and actual agent-tick denominators. Label the result “fixed-world broadcast population scaling,” not a pure causal effect of agent count.

6 E2 preregistration: RecruitmentArena-6

6.1 Purpose and information structure

E2 isolates recruitment before altering Alem. RecruitmentArena is a deterministic, pure-Python six-agent environment. Each episode contains two Warrior-, two Forager-, and two Miner-oriented private capability vectors and private task-specific costs. Public task cards contain a stable task ID, three-dimensional demand, required team size two or three, reward, and deadline.

The information boundary is:

- an agent knows its own true capabilities and costs;
- all agents see public task cards and delivered recruitment records;
- other agents’ true capabilities and costs remain private;
- published claims are public but need not equal private truth;
- true values are used only for terminal scoring, calibration, and the analysis-only oracle; and
- no oracle value enters an agent prompt or live allocation decision.

Episodes last 12 recruitment rounds. Control records arrive after a one-round delay. Each agent may emit at most one recruitment-control record per round, capped at 256 UTF-8 bytes. Ordinary conversation is disabled before a team locks and then delivered only to members of that task-bound team. An agent may belong to at most one active team. Membership ends on completion, cancellation, or deadline expiry.

6.2 Scenario families

Four paired scenario families isolate different allocation problems:

1. one feasible complementary task;
2. two disjoint tasks;
3. competing tasks requiring a scarce capability; and
4. oversubscription with more candidates than roster slots.

Formation-only episodes complete a task deterministically when the locked roster’s true capabilities satisfy its card. This separates formation quality from embodied execution. E2e later adds a five-round execution bridge.

6.3 Decentralized methods and controls

Open volunteer quorum. Agents publish applications. Every peer derives the same first-valid roster from the delayed public ledger. Selected members explicitly accept before the team locks.

Mutual nomination. Agents publish candidate rosters. A team activates only when every listed member has published the same reciprocal roster.

Peer Contract Net. A deterministic task sponsor is selected from ordinary embodied agents by a stable task-ID hash. The sponsor publishes a call for proposals, candidates bid, the sponsor awards a roster from visible bids, and every selected member accepts.

Controls are:

- one preformed all-six communication team;
- balanced fixed $3 + 3$ teams;
- no team and no ordinary communication;
- deterministic random allocation; and
- an offline oracle ceiling that never enters an agent prompt.

6.4 Versioned formation protocol

The planned TFP1 control plane includes task announcement, application, nomination, call for proposals, bid, award, accept, decline, lock, cancel, complete, and lease-expiry events. The canonical runtime must enforce:

- strict parsing and canonical serialization;
- the 256-byte UTF-8 limit by rejection, never truncation;
- one team per agent;
- reciprocal roster agreement before lock;
- task-bound completion, cancellation, and expiry;
- one-round delivery delay;
- public control versus team-private ordinary routes;
- deterministic transition sequence numbers and replay hashes; and
- zero ordinary cross-team deliveries.

The runtime validates events and routes. It must not choose an environment action or silently repair a failed model record.

6.5 E2a implementation amendment before canonical execution

A provider-free diagnostic matrix exposed three design ambiguities before any canonical E2 result was created. We therefore froze the following clarifications on 23 July 2026.

All task cards are active concurrently. A locked roster remains committed until a common allocation close after twelve acting rounds and a one-round delivery drain. Feasible teams are scored and completed together at that close. This prevents a scarce agent from being released and reused serially, and makes the live allocation comparable to the static `(tasks + idle)^6` oracle.

The weak Miner vector is (20, 20, 35). The provisional value 65 was invalid for the intended scarcity test: with additive team coverage, a nonspecialist’s 20 plus 65 already exceeds the mining demand of 80. Only the strong Miner at (20, 20, 90) can now satisfy either scarce card.

Each decentralized protocol is crossed with two bounded task-choice policies:

`local_commit.`

Each agent selects one task from public cards using only its own capability, its private cost, and a stable public tie-break. Mutual Nomination uses only the public role proxy because that protocol publishes no capability bid. Interest pools are disjoint and communication is low.

`public_sweep.`

Every agent submits at most one bounded record per round while sweeping all tasks in stable order.

Candidate overlaps are resolved from the delayed public ledger; no private cross-agent state enters the decision. This spends more control bytes to expose a larger candidate set.

The promotion denominator is oracle-allocation coverage: completed tasks divided by the number simultaneously allocable by the exact oracle. Individually feasible-card formation is reported separately; in the scarce family both cards are individually feasible even though at most one can be allocated. We additionally record achieved utility, utility regret only at equal reward, achieved and oracle raw cost, the oracle assignment, candidate-roster overlap, multi-offer rounds, and roster revisions.

The all-six and fixed 3 + 3 controls are topology references rather than exact rosters. Their raw reward and communication are retained, but normalized oracle reward and regret are marked not applicable. Fixed groups receive tasks by public order without consulting hidden capability truth, and no group is reused in the simultaneous allocation window.

6.6 E2b sequential selector amendment before hosted execution

The original E2b preregistration left a truth-free selector extension point but set both methods to `selector=none`. After that freeze, and before any hosted E2b response, the canonical provider-free E2d2 replication completed 12,000 episodes. Under truthful public claims, `joint_exact_allocation` achieved complete task coverage, matched the oracle allocation cost, and had zero utility regret in the constructed cross-task conflict. This result does not test strategic misreporting, model compliance, or embodied execution.

The prospectively dated 23 July 2026 amendment therefore changes only the live selector attached to Open Volunteer/Public Sweep:

- Open Volunteer replicas compute `joint_exact_allocation` from public task cards, delivered application self-claims, and current public leases.
- Mutual Nomination remains its native reciprocal rule, labelled `native_mutual_reciprocal`, as the cost-sensitive comparator.
- The pure selector receives no other agent’s true capability or cost, pending control, feasibility outcome, or oracle value.
- The enumeration assigns each available agent one task-or-idle label; every selected team therefore has the exact requested size and allocations are cross-task exclusive.
- The directory independently revalidates delivered applications, claimed feasibility, exact size, and public lease conflicts. Members still issue delayed `ACCEPT` controls and the native sponsor still issues `LOCK`.

This is a sequential mechanism substitution inside the frozen two-method matrix, not an added arm. The three seeds, four families, 24 episodes, 12-round horizon, model settings, output grammar, canary gate, and hard caps remain unchanged: 2,160 logical calls, 4,320 provider reservations, and 73,543,680 worst-case provider-attempt token reservations. Focused provider-free tests reproduce the E2d2 conflict,

show invariance to truth-field perturbations outside public claims, establish six-replica agreement, and exercise exact exclusive plans through lock and replay. No E2b behavioral result is yet available. An amended provider-free fake-client staging smoke passed the canary gate and resumed through all 24 atomic markers, using 416 fake logical calls and 832 provider-attempt reservations; these values validate orchestration only.

6.7 E2b prelaunch safety amendment before hosted execution

A second prospectively dated amendment was applied on 23 July 2026 after an independent prelaunch review and still before any hosted E2b response. It changes provenance, dispatch accounting, and response-envelope validation, but no prompt content, selector, scenario, behavioral outcome, or comparison. Hosted launch now fails closed before output creation or client construction unless credentials, canonical Git **HEAD** and status, **uv.lock**, all source hashes, this protocol, and the config hash agree. The only permitted dirty path is the exact untracked global replay artifact, whose content hash is recorded.

One nonblocking lock is held over the entire output root. Before each dispatch, a hash-chained JSONL ledger durably appends and **fsyncs** a reservation keyed by cell, round, agent, and semantic attempt. A crash leaves that reservation spent and stops automatic resume; completed-cell markers must exactly cover all reservations and resolutions. Cross-cell execution is sequential by default. Explicit **-parallel-cells -workers N** is safe because cells write disjoint paths while reservation append and all three campaign budgets are locked. Within-round six-agent concurrency is unchanged, and canary promotion remains sequential.

The input reservation now adds 1,024 framing tokens to the conservative one-token-per-prompt-byte charge. Consequently the executable cap is 77,967,360 provider-attempt tokens; the all-repair theoretical maxima are 62,373,888 recorded-success tokens and 124,747,776 provider-attempt tokens. Actual input above prompt bytes plus framing, output above 1,024, or more than two attempts fails closed, and provider **Retry-After** is clamped to 30 seconds.

For this arm only, the Responses adapter preserves the completion exactly, including outer whitespace and CR/LF. A valid envelope must be **completed**, have no incomplete reason, have a nonempty response ID, and return either the exact **gpt-5.6-luna** alias or that alias with a valid **YYYY-MM-DD** snapshot suffix. The resolved model is stable in the canary and exact-bound throughout the full stage. Marker schema v2 recomputes prompt bytes, raw hashes, the strict 256-byte TFP1 parse, normalized records, call/token counts, reservations, replay state and audit chains, deterministic hashes, and terminal analysis. The stored canary gate must be canonically identical to a pure recomputation over the one expected cell, and every gate predicate must be true.

Provider-free regressions cover crash reservations, concurrent-lock exclusion, ledger and marker tampering, copied cells, changed derived analysis, wrong or dated-invalid models, incomplete responses, missing IDs, raw whitespace/newlines, 256/257-byte boundaries, usage overages, and retry clamping. A post-hardening fake canary-to-full lifecycle used three explicit cell workers and completed all 24 markers with 416 logical calls, 832 provider reservations, 4,213,422 framing-aware token reservations, no unresolved reservation, and no overage. The first attempt exposed an overly global mid-run unresolved check; the corrected validator scopes that check to the cell while the final campaign gate remains global, and a regression pins this behavior. These checks establish orchestration integrity only; no E2b behavioral evidence has yet been collected.

6.8 E2b adversarial recovery amendment before hosted execution

A third prospectively dated amendment was applied on 23 July 2026 after a second independent adversarial review and before any hosted or network-backed E2b response. It changes no prompt, role, selector, seed, scenario, horizon, outcome, or cap. It closes five recovery paths by advancing the campaign, protocol/config, prompt-cache, and default output identities to v3 and the anchored reservation ledger to v2.

First, E2b now explicitly requests a strict Responses envelope. Missing model identity, missing usage, and missing or non-exact `input_tokens/output_tokens` fail before a call record can be accepted. Booleans, strings, floats, and nulls are never coerced. A present cache or reasoning counter is subject to the same exact-integer rule, while an omitted optional detail remains zero. The strict switch defaults off in the shared adapter, independently of exact completion preservation, so ordinary adapter behavior is unchanged. A malformed envelope consequently cannot write a comprehensively valid cell or promotion gate.

Second, reservation ledger v2 couples every reservation and resolution record to a separately created, exclusive, `fsynced` hash-chain anchor. Each manifest checkpoint binds record count, canonical ledger-prefix hash, record-chain head, anchor count, and anchor-chain head. The 18 logical calls of a complete canary therefore produce 36 anchored events; truncating a valid suffix while leaving the durable anchor high-water is a hard failure, not a fresh 18-call budget. Existing files are also fail-closed: a partial artifact triple, a complete but invalid triple, an unexpected managed file, or any cell file beside an empty ledger cannot be classified as pending.

Third, the managed root must be both lexical-canonical and realpath-contained. Managed opens use `O_NOFOLLOW`, regular-file checks, and link count one. Symlinked components and artifact directories, parent/root aliases, hardlinked locks, ledgers or artifacts, and two separately locked roots that share an artifact are rejected. Creation durably synchronizes every newly created ancestor, the directory itself, and its parent.

Finally, any provider-usage overage or reservation-resolution integrity failure permanently poisons the campaign. No new reservation may dispatch; queued cell futures are cancelled when possible, whereas an already in-flight request may only complete its durable resolution. The final bound manifest records the poison and reason. A failed or poisoned manifest cannot be resumed to reset that state, and ledger overages or unresolved reservations independently stop recovery.

The focused offline adversarial suite exercises exact usage and model identity, malformed-envelope non-promotion, valid-prefix canary truncation, partial/invalid/empty-ledger artifacts, path aliases and cross-root sharing, ancestor durability, sticky poison, in-flight resolution, resolution callback failure, and poisoned resume. A fresh fake-client lifecycle passed canary, then completed all 24 cells with three explicit cell workers: 464 logical calls, 928 provider-attempt reservations, 4,694,278 reserved tokens, 928 ledger events and 928 anchors, with zero unresolved reservations, zero overages, and no poison. Repeating full resume constructed no provider client and preserved the exact 928-event checkpoint. These measurements validate only offline orchestration and recovery. They provide no evidence yet about Luna recruitment behavior or comparative efficacy.

6.9 E2b v2 hosted failure and prospective v3 amendment

The first bounded hosted attempt was launched from commit `345da487a0217776b89b655e693dadcbfb153807` into the explicit v2 root. Its Open Volunteer/single-complementary canary passed all then-current gates. The promoted three-worker stage subsequently failed and poisoned itself: the full manifest contains 3 completed-marker entries, 2 failed Mutual cells, and 19 cancelled cells after 72 logical calls, 144 provider reservations, and 706,642 reserved tokens. This is a failed diagnostic, not an incomplete method comparison. The root is preserved and cannot resume.

The read-only summarizer binds this evidence in `Results/e2b_v2_failed_diagnostic_v1.json` and its Markdown rendering. The root-tree, full-manifest, and ledger SHA-256 values are, respectively, `406b4f09c1d5ae8e532bfaee0c9aa591324eb0603f0063635b54997c1ed374e0`, `14c90f60f5348472f00421842e552d623b282fa1a13ad08b4563d8a72d84fbd0`, and `950f5a7ce4ec3386bba2853e6409a5310496e5d5c69bc061a59e96cf3bcd4f04`. It classifies only the original 18-call Open/single canary as an eligible partial diagnostic. Open/two-disjoint made 22 calls and reached 0.5 oracle-allocation coverage, but it also records 12 poison-budget abstentions and is explicitly excluded. A zero-call poison artifact, both failed Mutual cells, and all cancelled cells are likewise excluded; no value is imputed.

The archived call envelopes identify a concrete output-budget failure. Nine of 72 calls returned `incomplete/max_output_tokens`; every one used exactly 1,024 output tokens, all 1,024 were reasoning tokens, and the visible completion was blank. The single and two-disjoint Mutual cells contained four and five such calls, respectively. Two other Mutual outputs were well-formed nominations that omitted their own sender from the roster and failed `semantic.sender_missing`. Thus the evidence supports an output-censoring diagnosis and records a separate coordination error, but does not support a comparative efficacy claim.

Because completed responses already reached 1,020 output tokens, the right-censored sample cannot establish that 2,048 would clear the high-reasoning tail. Before any v3 response, the successor therefore freezes 4,096 output tokens. This is four times the censoring boundary, while the input-dominated per-attempt reservation grows only 17.0%, from 18,048 to 21,120 tokens. The visible TFP1 grammar remains capped at 256 bytes. Any v3 max-output truncation is still invalid; the cap is never adapted after a v3 observation.

V3 also replaces the one-method canary with four seed-22000 cells: both Open Volunteer and Mutual Nomination crossed with single-complementary and two-disjoint. Each cell must independently pass replay, envelope, transport, invalid-call, repair, budget, and no-truncation gates and form at least one truly feasible team. All requested/resolved model bindings must agree. A missing, failed, invalid, poison-affected, zero-call, or nonforming cell blocks full-stage client construction.

The full logical/provider limits remain 2,160 and 4,320; the 4,096-token output allowance raises the hard provider-attempt exposure to 91,238,400. The four-cell canary projects at most 360 logical calls, 720 provider reservations, and 15,206,400 reserved tokens under the 25% repair gate. A fresh offline fake lifecycle passed all four canary cells and completed 24/24 full cells with three workers: 440 logical calls, 880 provider reservations, 7,157,450 reserved tokens, 880 ledger events/anchors, no unresolved reservation, no overage, and no poison. A repeat full resume constructed no client and changed no checkpoint. No hosted v3 call has been made.

6.10 Member-selection ablation

Holding Peer Contract Net fixed, E2d compares:

1. `first_valid`: first claimed-feasible roster by arrival round and agent ID;
2. `random_valid`: deterministic seeded sampling from claimed-feasible rosters; and
3. `exact_utility`: exact enumeration of subsets of the required size.

For claimed capability vector q_i , cost c_{it} , and task demand d_t , the exact rule maximizes

$$U(S, t) = \min_{j: d_{tj} > 0} \frac{\sum_{i \in S} q_{ij}}{d_{tj}} - 0.25 \frac{\sum_{i \in S} c_{it}}{100|S|}.$$

A roster is claimed-feasible only if it has the requested size and coverage of at least one in every demanded dimension. Utility ties resolve by lexicographically sorted agent IDs. The analysis-only oracle enumerates (tasks + idle)⁶ assignments using true capabilities and costs.

6.11 E2 stages and hypotheses

H2a: formation.

At least one decentralized method will form feasible teams before deadline in at least 80% of feasible opportunities.

H2b: structure.

Contract Net will reduce conflicting rosters relative to open volunteer formation, at the cost of additional control latency.

H2c: selection.

Exact utility will reduce oracle regret relative to first-valid and random selection when scarce capabilities create competing feasible rosters.

Stage	Seeds	Model use	Purpose
E2a	20000–20999	None	Scripted correctness and determinism
E2b	22000–22002	GPT-5.6 Luna high	LLM mechanism screen
E2c	22100–22109	GPT-5.4 nano high	Paired efficacy estimate
E2d	22200–22209	Fixed contract net	Selection-rule ablation
E2e	22300–22309	GPT-5.4 nano high	Five-round execution bridge

Table 5: E2 funnel. Only methods passing the mechanism gate enter E2c.

H2d: communication boundary.

Team-private execution will use fewer delivered ordinary-message bytes than all-six communication while maintaining complementary completion.

Primary E2 outcomes are normalized task reward and oracle regret. Secondary outcomes are formation latency, feasible-team rate, churn, roster agreement, truthful-report calibration, messages, delivered bytes, tokens, and wall time.

Planned figure: Paired normalized reward and oracle regret by recruitment method, followed by formation latency and delivered-byte efficiency.

Pending result. Record E2a as an engineering result before any paid call. For E2b–E2e, preserve raw model output and rejection codes; do not rewrite invalid messages. If no method passes the mechanism gate, stop before E3, version the protocol, and preserve the negative result.

7 E3 preregistration: recruitment in Alem

7.1 Treatment semantics

E3 introduces recruitment only through opt-in topologies. The baseline topology remains behaviorally unchanged.

- An unteamed agent may emit only one valid bounded TFP1 record on the public bulletin.
- Invalid, oversized, or free-form unteamed messages are rejected, logged, and delivered to no one.
- Once a roster locks, ordinary Source-style `<communication>` text is delivered only to current teammates.
- Nonmembers receive no ordinary team-message copy.
- Public control traffic and private ordinary traffic are charged and reported separately.
- Membership changes take effect deterministically after the accepted lock transition and are replayable from the event journal.
- Middleware validates and routes; it does not select an Alem action.

This is a team-scoped *communication* study, not strict information isolation. Ordinary Alem teammate observations remain visible. Environment Requests and targeted **Give** are not filtered by team in the initial E3 study.

7.2 Observable tasks and stable identity

An embodied agent must locally observe an opportunity before announcing it. Stable task IDs derive only from observable task type, dungeon level, and location. The current text wrapper is insufficient by itself because it collapses multiple same-type coordination cues and normally omits absolute coordinates.

E3 therefore requires structured, locally visible opportunity metadata generated from the same light and view masks as the agent’s text. This metadata may validate and identify an agent-selected announcement; it must not trigger an automatic announcement or expose a hidden opportunity.

Before E3a, the task taxonomy and capability mapping must be frozen for coordinated mining, construction, elite combat, and handover opportunities. All arms in the first E3 study use the 48-by-48 world and task requirements of two or three agents.

7.3 Critical environment control

At six agents, current world generation sometimes assigns the full physical population as the requirement for mining, construction, and elite-mob coordination. Although an environment field named `max_agents_required` exists, current generation does not apply it. Changing this globally would silently alter the Source scaling environment.

The E3 implementation must therefore add an explicit, experiment-only two-to-three-agent requirement control. It must be enabled identically for every E3 comparison arm and disabled by default so E0/E1 world generation is unchanged.

7.4 Arms, stages, and hypotheses

E3 compares:

1. unchanged six-agent Source full broadcast;
2. balanced fixed 3 + 3 communication teams;
3. dynamic recruitment using the promoted E2 method and selection rule;
4. a bulletin-only, no-ordinary-conversation screening floor.

The final paper-scale comparison retains Source, fixed 3 + 3, and dynamic recruitment. A second recruitment method is retained only if its normalized reward is within 5% of the best method while materially reducing communication.

Stage	Environment	Seeds	Horizon	Purpose
E3a	Alem Debug, scripted god mode	14000–14001	20	Route correctness
E3b	Alem Debug, Easy	14100–14102	200	Paid mechanism screen
E3c	Full Alem, Easy	14200–14209	1,000	Paired efficacy estimate
E3d	Full Alem, all difficulties	14300–14319	10,000	Paper scale

Table 6: E3 funnel. E3d requires explicit cost authorization.

H3a: transfer.

A mechanism passing E2 will form valid task teams under Alem’s observation and survival constraints.

H3b: communication efficiency.

Dynamic recruitment will reduce ordinary delivered bytes relative to all-six Source broadcast.

H3c: performance preservation.

The communication reduction will not reduce any co-primary outcome by more than the preregistered 5% relative margin in the final paired comparison.

H3d: adaptivity.

Dynamic teams will outperform fixed 3 + 3 teams on scenarios where task locations and capability needs change during the episode.

Planned figure: Performance–communication frontier for Source broadcast, fixed 3 + 3, dynamic recruitment, and the bulletin-only floor.

Pending result. E3 results may be added only after E3a demonstrates zero unauthorized ordinary deliveries, deterministic membership replay, and unchanged baseline prompts/routes. Report locally observed task provenance and every rejected control record.

8 Promotion rules and statistical reporting

8.1 Mechanism gates

Every implemented recruitment mechanism must satisfy:

- zero unauthorized cross-team ordinary deliveries;
- 100% roster agreement for every activated team;
- at least 95% action and recruitment-transition validity;
- zero unrecovered provider failures in canonical episodes;
- at least 80% of feasible opportunities formed before deadline;
- matching deterministic replay hashes; and
- no baseline prompt or route change under baseline topology.

These are eligibility gates, not outcomes that may be averaged away. A routing violation blocks efficacy interpretation even if reward is high.

8.2 Behavioral promotion

For a ten-seed behavioral promotion, a treatment must have positive paired direction on at least seven of ten seeds for at least two of Total percentage, mean per-agent return, and unique team achievements, with no material regression in the third. Effect sizes and paired uncertainty are reported even when this directional rule fails.

A final communication-efficiency claim requires a paired 95% interval showing no worse than a 5% relative performance margin on all three co-primary outcomes while reducing delivered communication bytes by at least 30%. Communication cost and model-token cost are reported separately.

8.3 Multiple views of cost

Every behavioral table includes:

- total model calls and calls per actual agent-tick;
- input, output, reasoning, cache-read, and cache-write tokens;
- emitted payload bytes and fan-out-weighted delivered bytes;
- control-plane and ordinary-message bytes separately;
- summed model latency and complete tick wall time; and
- requested versus executed ticks and alive-agent-step exposure.

The coordination-efficiency diagnostic is

$$\text{coordination efficiency} = \frac{\text{paired team-performance improvement}}{\text{delivered communication bytes}},$$

reported alongside its numerator and denominator rather than as a standalone ranking.

9 Implementation readiness and risk register

E2 now has a deterministic TFP1 runtime, RecruitmentArena, selectors, analysis-only oracle, provider-free tests, and a staged resumable E2b launcher. Its fake-client canary-to-full smoke validates orchestration and accounting, not behavior. E3-specific Alem task extraction and opt-in receiver-specific routing remain absent. The evaluator supplies parallel calls, one-tick broadcast, route accounting, and canonical artifacts, while E0 supplies six-agent compatibility evidence. The older DCP1 parser remains distinct from TFP1 and must not be used as evidence for task cards, membership, leases, byte caps, or receiver-specific routing.

Table 7: Pre-implementation E2/E3 risk register.

Risk	Failure mode	Required mitigation
Baseline route drift	A routing refactor changes Source delivery, order, delay, or accounting.	Keep the existing baseline branch explicit; add golden prompt and $N(N - 1)$ route regressions before dynamic routing.
Global task-size change	Applying a three-agent cap globally changes E1’s $N = 4/6$ worlds.	Use an opt-in E3-only requirement control, identical across E3 arms and disabled by default.
Unstable task IDs	Text cues omit absolute location or collapse multiple opportunities.	Expose structured locally visible opportunities from the same visibility mask; never inspect or announce hidden tasks.
Byte-limit corruption	Truncation turns one invalid record into a different apparent record.	Measure the complete UTF-8 payload and reject records above 256 bytes.
Hidden allocation	Middleware or oracle chooses a roster or action using private truth.	Use only published claims for live selection; reserve true values for scoring and the offline oracle.
Roster disagreement	Agents act on different membership states after delayed messages.	Require reciprocal acceptance, canonical transition order, per-agent ledger comparison, and deterministic replay.
Cross-team leakage	One sender-indexed broadcast map delivers private text to all agents.	Use receiver-indexed delivery envelopes for opt-in team topologies and assert zero nonmember recipients.
Protocol noncompliance	LLM output is malformed, oversized, stale, or unauthorized.	Log stable rejection codes, deliver nothing, and never silently rewrite a live record.
Prompt confound	Recruitment instructions accidentally alter the baseline condition.	Inject formation context only under explicit recruitment/fixed-team settings; snapshot baseline prompts.
Premature efficacy claim	Reward is analyzed despite mechanism or replay failure.	Apply structural promotion gates first and pre-serve stopped/negative studies.
Cost confound	Different episode-worker concurrency changes provider load by arm.	Use one episode worker, concurrent within-tick agents, randomized sequential blocks, and complete usage accounting.

10 Audit and provenance requirements

Every canonical study records:

- resolved configuration and exact CLI invocation;
- Git commit, branch, dirty-source manifest, dependency lock/hash, Python version, and platform;
- model snapshot, reasoning effort, cache policy, and nominal call ceiling;
- scenario or environment seeds and paired-arm mapping;
- raw prompts, raw responses, provider response IDs, finish reasons, usage, retries, and transport errors;
- parsed actions and recruitment events with stable validation codes;
- task cards, event sequence, roster transitions, leases, route envelopes, rejected deliveries, and replay hashes;
- terminal metrics and actual agent-tick denominators; and
- the exact CSV/JSON inputs for every table and figure.

Artifacts from failed or partial attempts are archived rather than overwritten. Canonical episodes require terminal status and no API/runtime error. Resumption retries only failed seeds. Paid stages begin with a dry-run call/token ceiling and a clean committed implementation.

10.1 Completed E0 provenance

E0 ran from 2026-07-23T06:24:16.409143Z to 2026-07-23T06:28:23.047401Z at commit 587e7176ce3b02296120e9504c1d7d609e9dcf2a. Its raw root is `outputs/alem_eval/e0_agent_compatibility_v1`. The source-of-truth result SHA-256 is 7b67b56544a00df63b1e07106c9eeffe51eedc9bc62dc7d8d914af734ad8c47d.

10.2 Completed E0.1 provenance

The complete E0.1 matrix ran from 2026-07-23T06:57:25.963656Z to 2026-07-23T07:08:25.236294Z using measured-run commit 5c4c64338b9aae9132f9cc32b6eec3d746964529. Its raw root is `outputs/alem_eval/e0_1_agent_compatibility_200_v2`. The source-of-truth result SHA-256 is 6d28e885e241b4f58553291d6fa205fb982878f2575f6bdd3a24d711fda9b529. The earlier diagnostic attempt remains at `outputs/alem_eval/e0_1_agent_compatibility_200_v1`; it stopped after the two $N = 1$ episodes under the original strict validator.

Both source manifests recorded the pre-existing untracked replay `Results/replays/nano_high_source_full_world.mp4`. It was outside the run roots and was not an input to either result.

11 Interpretation rules and limitations

- E0 and E0.1 establish engineering compatibility only. Their Noop returns and mortality are not behavioral scaling evidence.
- E1 is a descriptive fixed-world population curve. It cannot isolate agent count from role balance, spawn geometry, prompt growth, mob pressure, or coordination requirements.
- $N = 1$ uses a different solo wrapper, and $N = 6$ is beyond the paper’s documented one-to-four range.
- E2’s formation-only reward is a mechanism diagnostic, not evidence that a team can navigate or execute in Alem. E2e and E3 address that gap.
- E2b v2 is a failed, method-incomplete campaign. Its one eligible Open canary and excluded poison-affected artifacts cannot estimate an Open-versus- Mutual effect. The 4,096-token v3 allowance is a prospective response to right-censoring, not evidence that truncation has been eliminated; the four-cell canary must test that claim.
- E2b’s rollback defense creates and synchronizes one anchor for every reservation and resolution: two anchor files per logical provider call. This does not add model calls or tokens, but it adds local metadata

and `fsync` latency. The fake 464-call lifecycle created 928 anchors; hosted wall-time overhead has not yet been measured and must be reported separately from provider latency.

- E3 isolates ordinary message routing, not all channels of information or resource transfer.
- High reward cannot validate a method that leaks messages, disagrees on rosters, or uses hidden information.
- A result from three seeds is a mechanism screen, not confirmatory efficacy evidence. Ten- and twenty-seed stages retain raw paired effects and uncertainty.
- E3d and other 10,000-tick paid stages require explicit authorization after a fresh estimate.

12 Planned final synthesis

The final report will answer the research questions in causal order:

1. Does unchanged Source behavior improve, plateau, or degrade as N increases, and at what communication/token cost?
2. Which decentralized recruitment protocol most reliably forms feasible teams under the bounded bulletin?
3. Does mathematically informed selection reduce oracle regret?
4. Does the promoted mechanism preserve performance when execution and survival are restored?
5. Does team-scoped communication reduce delivered bytes without exceeding the 5% performance-regression margin?

Pending result. Replace this box with a claim-by-evidence table after E1–E3. Every final claim must cite a completed stage, paired seed set, uncertainty statement, mechanism-gate result, and cost outcome. Claims unsupported by a completed stage remain explicitly prospective.

13 Current conclusion

The only completed conclusion is engineering. Alem’s canonical language-agent evaluation path supports one, two, three, four, and six physical agents under the tested provider-free conditions. E0 passed its short compatibility gate. E0.1 passed every structural gate but failed the literal full-horizon requirement because three inactive-policy episodes ended naturally.

The behavioral and decentralized-coordination questions remain open. The preregistered E1–E3 sequence is designed to prevent a routing, formatting, hidden-information, or population confound from being mistaken for collective intelligence. This report must be updated from canonical artifacts rather than from console summaries as each downstream gate is completed.

A Reproduction commands for completed gates

E0.

```
uv run -extra baselines-llm -python 3.12 \  
python scripts/run_e0_agent_compatibility.py \  
-output outputs/alem_eval/e0_agent_compatibility_v1
```

E0.1.

```
uv run -extra baselines-llm -python 3.12 \  
python scripts/run_e0_agent_compatibility.py \  
-counts 1,2,3,4,6 \  
-seeds 13000,13001 \  
-steps 200 \  
-output outputs/alem_eval/e0_1_agent_compatibility_200_v2
```

B Planned result-table inventory

1. E1 population-by-seed co-primary outcomes and cost metrics.
2. E1 paired bootstrap intervals and population trend contrasts.
3. E2a scenario-by-method mechanism gates.
4. E2b/c raw validity, formation, reward, regret, and cost.
5. E2d roster-selection utility and oracle-regret contrasts.
6. E2e formation-to-execution attrition and team-private delivery audit.
7. E3a route-invariant and task-provenance audit.
8. E3b/c paired Source, fixed-team, and dynamic-team outcomes.
9. E3 performance–communication noninferiority analysis.
10. Final claim-by-evidence and negative-result table.

C Deferred extensions

Strict filtering of teammate observations, Requests, and Give; message loss; crash-stop and Byzantine agents; reputation learning; teams larger than three; populations above six; learned routing; and long-horizon context or role-drift studies are outside E0–E3. They require separate protocols rather than post-hoc interpretation of this program.